


FUNDAMENTALS OF DEEP LEARNING

Part 2: How a Neural Network Trains



AGENDA

Part 1: An Introduction to Deep Learning

Part 2: How a Neural Network Trains

Part 3: Convolutional Neural Networks

Part 4: Data Augmentation and Deployment

Part 5: Pre-trained Models

Part 6: Advanced Architectures

AGENDA – PART 2

- Recap
- A Simpler Model
- From Neuron to Network
- Activation Functions
- Overfitting
- From Neuron to Classification

RECAP OF THE EXERCISE

What just happened?

Loaded and visualized our data

Edited our data (reshaped, normalized, to categorical)


Created our model

Compiled our model

Trained the model on our data


DATA PREPARATION

Input as an array




DATA PREPARATION

Targets as categories





AN UNTRAINED MODEL



A SIMPLER MODEL

$$y = mx + b$$



x	y
1	3
2	5



A SIMPLER MODEL

$$y = mx + b$$



x	y
1	3
2	5



A SIMPLER MODEL

$$y = mx + b$$


x	y	\hat{y}
1	3	4
2	5	3



A SIMPLER MODEL

$$y = mx + b$$

x	y	\hat{y}	err^2
1	3	4	1
2	5	3	4
MSE =			2.5
RMSE =			1.6




$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

A SIMPLER MODEL

$$y = mx + b$$

x	y	\hat{y}	err^2
1	3	4	1
2	5	3	4
MSE =			2.5
RMSE =			1.6





```



1 data = [(1, 3), (2, 5)]
2 m = -1
3 b = 5
4
5
6 def get_rmse(data, m, b):
7     """Calculates Mean Square Error"""
8     n = len(data)
9     squared_error = 0
10    for x, y in data:
11        # Find predicted y
12        y_hat = m*x+b
13        # Square difference between
14        # prediction and true value
15        squared_error += (
16            y - y_hat)**2
17    # Get average squared difference
18    mse = squared_error / n
19    # Square root for original units
20    return mse **.5
21

```



THE LOSS CURVE



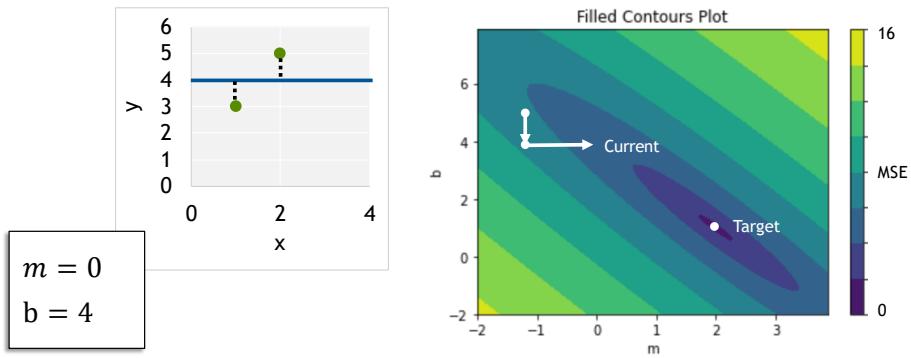
THE LOSS CURVE




THE LOSS CURVE



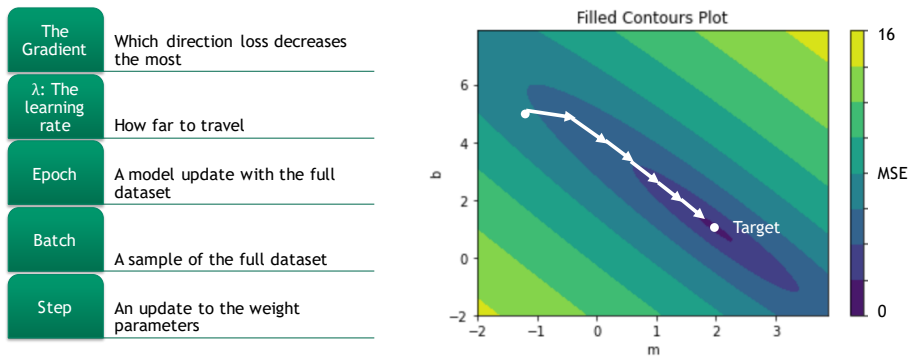
THE LOSS CURVE




THE LOSS CURVE



THE LOSS CURVE




OPTIMIZERS



- Adam
- Adagrad
- RMSprop
- SGD




BUILDING A NETWORK




- Scales to more inputs

BUILDING A NETWORK



- Scales to more inputs
- Can chain neurons

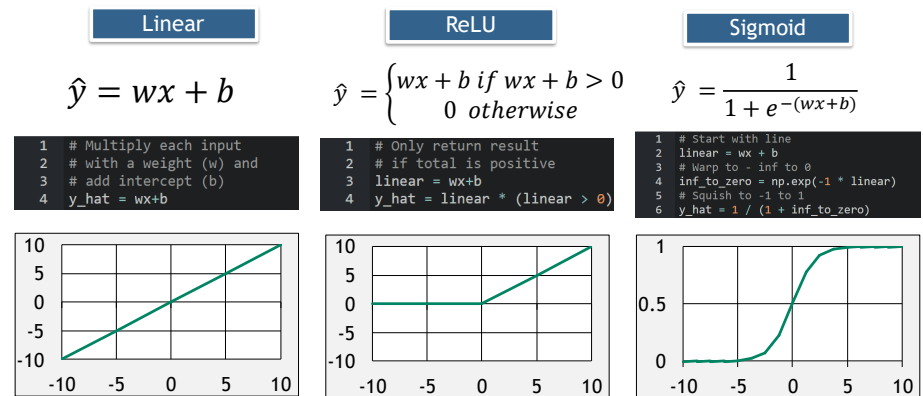
BUILDING A NETWORK



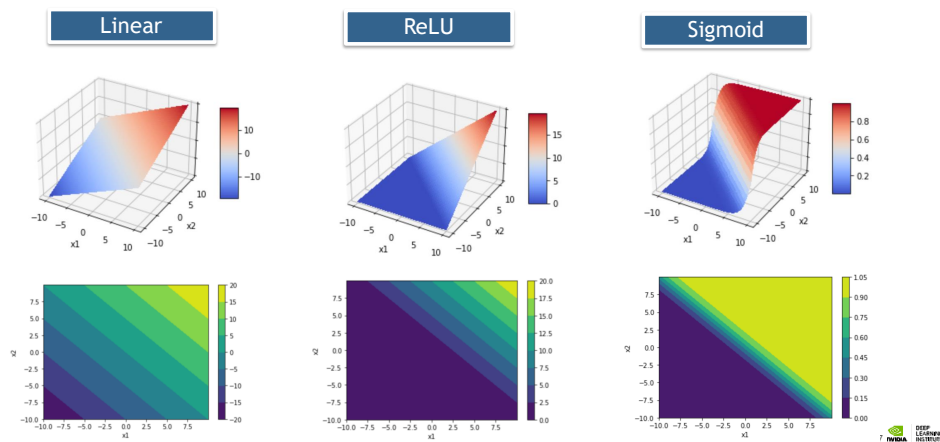
- Scales to more inputs
- Can chain neurons
- If all regressions are linear, then output will also be a linear regression



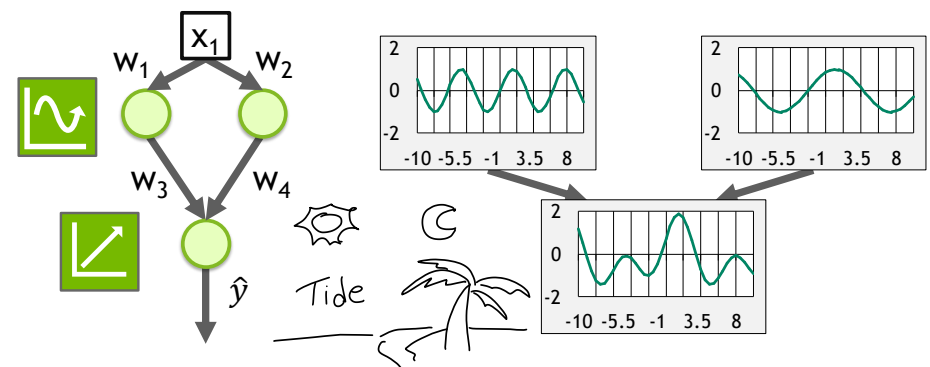
ACTIVATION FUNCTIONS



ACTIVATION FUNCTIONS



ACTIVATION FUNCTIONS







OVERFITTING

Why not have a super large neural network?





OVERFITTING

Which Trendline is Better?



OVERFITTING

Which Trendline is Better?



TRAINING VS VALIDATION DATA

Avoid memorization

Training data


- Core dataset for the model to learn on

Validation data

- New data for model to see if it truly understands (can generalize)

Overfitting


- When model performs well on the training data, but not the validation data (evidence of memorization)
- Ideally the accuracy and loss should be similar between both datasets



33 DEEP LEARNING INSTITUTE




AN MNIST MODEL




35 DEEP LEARNING INSTITUTE

AN MNIST MODEL




36 DEEP LEARNING INSTITUTE


AN MNIST MODEL





RMSE FOR PROBABILITIES?




RMSE FOR PROBABILITIES?



CROSS ENTROPY




CROSS ENTROPY




$$\text{Loss} = -((t(x) \cdot \log(p(x)) + (1 - t(x)) \cdot \log(1 - p(x)))$$

$t(x)$ = target (0 if False, 1 if True)


$p(x)$ = probability prediction of point x



CROSS ENTROPY




```
1 def cross_entropy(y_hat, y_actual):
2     """Infinite error for misplaced confidence."""
3     loss = log(y_hat) if y_actual else log(1-y_hat)
4     return -1*loss
```



BRINGING IT TOGETHER


THE NEXT EXERCISE

The American Sign Language Alphabet





DEEP
LEARNING
INSTITUTE



LET'S GO!

APPENDIX: GRADIENT DESCENT

HELPING THE COMPUTER CHEAT CALCULUS

Learning From Error

$$MSE = \frac{1}{n} \sum_{i=1}^n (y - \hat{y})^2 = \frac{1}{n} \sum_{i=1}^n (y - (mx + b))^2$$

$$MSE = \frac{1}{2} ((3 - (m(1) + b))^2 + (5 - (m(2) + b))^2)$$

$$\frac{\partial MSE}{\partial m} = 5m + 3b - 13$$


$$\frac{\partial MSE}{\partial b} = 3m + 2b - 8$$

$$\frac{\partial MSE}{\partial m} = -3$$

$$\frac{\partial MSE}{\partial b} = -1$$

$m = -1$
 $b = 5$

THE LOSS CURVE



THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \quad \frac{\partial MSE}{\partial b} = -3$$



THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \quad \frac{\partial MSE}{\partial b} = -3$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$



THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \quad \frac{\partial MSE}{\partial b} = -3$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$

$$\lambda = .6$$



THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \quad \frac{\partial MSE}{\partial b} = -3$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$

$$\lambda = .005$$



THE LOSS CURVE

$$\lambda = .1$$

$$m := -1 + 7\lambda = -0.3$$

$$b := 5 + 3\lambda = 4.7$$

